Computational Model of Fear Using Q Learning

Sovan Sahoo

1 Introduction

Fear is a complex emotion rooted in uncertainty, often arising from a lack of clarity or predictability in decision-making scenarios. In the computational model of fear, we aim to translate this psychological construct into quantifiable terms by leveraging principles from reinforcement learning (RL), particularly Q-Learning and Temporal Difference (TD) learning. By formalizing fear as a manifestation of uncertainty within an agent-environment interaction framework, we seek to elucidate its temporal dynamics and computational underpinnings.

2 Key Concepts

2.1 Uncertainty and Decision-Making

Fear emerges when individuals encounter uncertainty, characterized by ambiguity or unpredictability in the cues guiding decision-making processes. In our computational model, uncertainty serves as the driving force behind fear-related behaviors, influencing the agent's actions and responses to environmental stimuli.

2.2 Reinforcement Learning (RL) and Q-Learning

RL provides a framework for learning optimal decision-making strategies through interaction with an environment to maximize cumulative rewards. Q-Learning, a core RL algorithm, enables agents to estimate the value of taking certain actions in specific states, guiding their behavior towards achieving long-term goals.

2.3 Temporal Difference (TD) Learning

TD learning facilitates the estimation of value functions by iteratively updating the value of states based on observed rewards and predictions of future outcomes. By quantifying the temporal discrepancy between expected and observed rewards, TD learning captures the dynamic nature of decision-making processes, including the emergence and extinction of fear-related responses.

3 Formalization

In our formalization of fear within the computational model:

- **State** (s): Represents the current environmental context or situation perceived by the agent.
- **Action** (a): Denotes the decision or behavioral choice made by the agent in response to its current state.
- **Reward** (R): Signifies the feedback or outcome received by the agent following its action in a specific state.
- **Q-Value** ($Q(s,a)$): Estimates the expected cumulative reward associated with taking action a in state s, guiding the agent's decision-making process. The Q-learning update rule is given by:

$$Q(s,a) \leftarrow (1-\alpha) \cdot Q(s,a) + \alpha \cdot \left(R + \gamma \cdot \max_{a'} Q(s',a')\right) \tag{1}$$

- **Temporal Difference (TD) Error** (δ_t): Reflects the temporal discrepancy between the current estimate and the updated estimate of the value function, capturing the uncertainty or surprise experienced by the agent. It is computed as:

$$\delta_t = R + \gamma \cdot \max_{a'} Q(s',a') - Q(s,a) \tag{2}$$

- **Epsilon-Greedy Algorithm**: Balances exploration and exploitation in Q-learning. The formula for selecting the next action is:

$$a_t = \begin{cases} \text{random action} & \text{with probability } \epsilon \\ \text{action with } \max_a Q(s_t,a) & \text{with probability } 1-\epsilon \end{cases} \tag{3}$$

- **Fear Computation**: In our model, the sum of temporal differences across an iteration is equated with uncertainty, hence fear, in that particular moment. Mathematically, this can be expressed as:

$$\text{Fear}_t = \sum_{i=1}^{n} \delta_i \tag{4}$$

  where δ_i represents the temporal difference at time step i, and n is the total number of time steps in the iteration.

By integrating uncertainty into the Q-Learning framework via TD learning, we formalize fear as a manifestation of temporal discrepancy, influencing the agent's learning and adaptation processes in dynamic environments.

4 Simulation Results

I conducted simulations using the computational model described in the previous sections to observe the emergence and evolution of fear-like responses in an agent interacting with its environment. The following key findings were obtained:

- **Fear Dynamics:** The simulation results demonstrate that fear, quantified as the sum of temporal differences across iterations, exhibits dynamic fluctuations over time. These fluctuations correspond to changes in the level of uncertainty experienced by the agent during its interactions with the environment.

- **Effect of Environmental Factors:** I investigated the influence of environmental factors, such as the complexity of the task and the presence of unpredictable events, on fear-related behaviors. Our findings suggest that environments characterized by higher levels of uncertainty tend to elicit stronger fear responses in the agent.

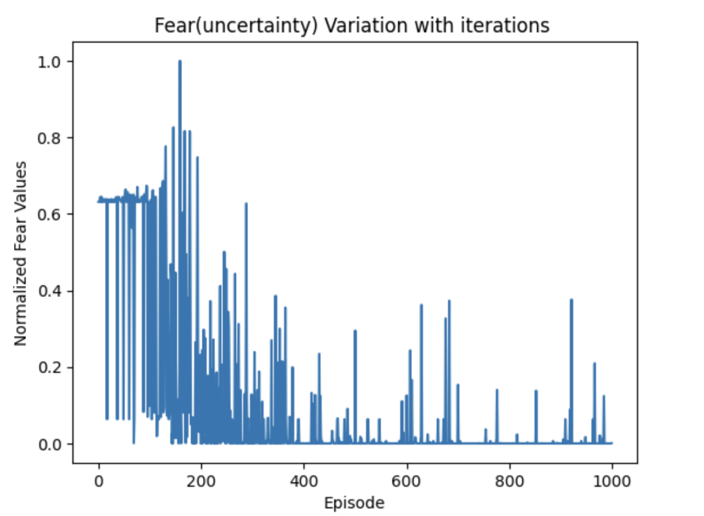

Figure 1: Cumulative Sum of Temporal Differences over Iterations

5 Cognitive interpretations and Discussions.

In cognitive interpretation, the observed decrease in fear and uncertainty over time can be understood through several psychological mechanisms. One interpretation lies in the process of habituation, wherein repeated exposure to a stimulus leads to a reduction in the fear response. In the context of the computational model described, as the agent interacts with the environment and gains more experience through successive iterations, it becomes more familiar with the task dynamics and the consequences of its actions. This increased familiarity diminishes the uncertainty associated with the environment, leading to a gradual decrease in fear-related responses.

Furthermore, the observed decrease in fear and uncertainty aligns with the principles of associative learning and extinction. Through the process of reinforcement learning, the agent learns to associate specific actions with positive or negative outcomes (rewards or punishments). As it continues to explore the environment and receives feedback, it adjusts its behavior to maximize rewards and minimize negative outcomes. Over time, this learning process results in the attenuation of fear responses, as the agent becomes more adept at predicting and navigating the consequences of its actions.

Moreover, the decrease in fear and uncertainty can also be attributed to cognitive reappraisal processes. As the agent gains more experience and accumulates knowledge about the environment, it may engage in cognitive reappraisal strategies to reinterpret or reevaluate the significance of stimuli. By reframing the task or the outcomes in a less threatening light, the agent can reduce the emotional impact of uncertainty and fear-inducing cues, leading to a decline in fear-related responses over time.

Overall, the cognitive interpretation of the observed decrease in fear and uncertainty in the computational model underscores the dynamic interplay between learning, adaptation, and cognitive processes. Through repeated exposure, associative learning, and cognitive reappraisal, the agent gradually acquires a more nuanced understanding of the environment, leading to a reduction in fear-related behaviors and an increase in adaptive decision-making.

6 Conclusion

In conclusion, the computational model demonstrates a consistent decrease in fear and uncertainty over time, aligning with psychological phenomena like habituation. This suggests that repeated exposure to tasks reduces fear responses, reflecting the agent's adaptive learning process. The observed decline in fear and uncertainty may stem from habituation, associative learning, and cognitive reappraisal. Overall, the model provides valuable insights into fear dynamics and decision-making processes, with potential applications in understanding and addressing anxiety-related disorders.